


Improving System Performance

Computer Engineering 1

Agenda

- Motivation
- Aspects of Optimization
- Bus Architectures
- Instruction Set Architectures (ISA)
- Pipelining
- Parallel Computing




Motivation

- **Intel Pentium E5800**

- 3.2 GHz
- 2 Cores
- 2 MB cache
- Released Q4 2010
- **PassMark: 1184**
- **TDP: 65 W**

spEED!




- **Intel Celeron G4930**

- 3.2 GHz
- 2 Cores
- 2 MB cache
- Q2 2019
- **PassMark: 2628**
- TDP: 54 W

- **AMD Ryzen r1102G**


- 1.2 GHz
- 2 Cores
- 4 MB cache
- Q1 2020
- PassMark: ~1600
- **TDP: 6 W**

LOW POWER!



Motivation

– Growth in processor performance




Source: John L. Hennessy, David A. Patterson (2017). Computer Architecture: A Quantitative Approach. 6th edition.

Learning Objectives

- **At the end of this lesson you will be able**
 - to explain different types of bus architectures
 - to understand the difference between von Neumann and Harvard architecture
 - to understand RISC and CISC paradigms
 - to describe the idea of pipelining
 - to calculate processing performance improvement through pipelining
 - to describe the basics of parallel computing


Aspects of Optimization

– ‘Wishlist of Optimization’



- This lecture: Increasing system/CPU speed
- Some other topics: Covered in ‘Microcomputer Systems 1’ (MC1)

How to Increase System speed?



Bus Architecture

How is the connection between the CPU and memory organized?


A single system bus: **simple, cheaper** **slower**

Multiple system buses: **faster** **more complexity**

Bus Architecture


- **von Neumann Architecture**

- Same memory holds program and data
- Single bus system between CPU and memory




- **Harvard Architecture**

- “Mark I” at Harvard University (Howard Aiken, 1939-44)
- Separate memories for program and data
- Two sets of address/data buses between CPU and memory




Bus Architecture

- von Neumann Architecture
 - Example: ARM Cortex-M0



Bus Architecture

- **Harvard Architecture**
 - Example: ARM Cortex-M3/M4



Instruction Set Architecture

Powerful, complex instructions?

or

Simple, highly optimizable instructions?

(or something in between?)

Instruction Set Architecture

- **First CPUs**

- Small address bus, small data bus (8 bit)
- Few registers, no cache, pipelining, co-processors, multi-core
- Machine code as finite state machine: Shift, multiply, etc. in loops

- **Consequences for machine code**

- Differing complexity
 - Different length
 - Variable execution time
 - Individual addressing / arguments
- 
- for different instructions

→ **More and more complex instructions (CISC)**

Instruction Set Architecture

- **John Cocke (IBM, 1974) proved that**
 - 80% of work was done using only 20% of the instructions
 - Most (complex) instructions were not used at all by compilers
- **RISC (Reduced Instruction Set Computer)**
 - Few instructions, unique instruction format
 - Fast decoding, simple addressing
 - less hardware → allows higher clock rates
 - more chip space for registers (up to 256!)
 - Load-store architecture reduces memory accesses, CPU works at full-speed on registers

Remark: Word “CISC” was introduced with “RISC”



Acorn Archimedes A3000
Source: Wikipedia Commons

Instruction Set Architecture

- **RISC**

- Load / Store Architecture
- Data processing instructions only available on registers

- **CISC**

- One of the operands of an instruction may directly be a memory location


Example: **Balance = Balance + Credit**

```
ldr    r0,=Credit
ldr    r1,[r0]
ldr    r0,=Balance
ldr    r3,[r0]
adds   r2,r1,r3
str    r2,[r0]
```

```
mov    AX, [Credit]
add    [Balance], AX
```

Instruction Set Architecture

– Instruction / opcode complexity




No RISC – No Fun

- **RISC advantages**
 - Simple and fast instruction decoding
 - More registers/cache on silicon area (less memory access)
 - Several data paths possible (superscalar computers)
 - Higher clock frequencies
 - Effective compiler optimization with limited, generic instructions
 - Easy pipelining, shorter pipelines (instruction size / duration)
- **CISC advantages**
 - Less program memory needed with complex instructions
 - Short programs may work faster with less memory accesses
- **CISC and RISC more and more indistinguishable**
 - Modern x86 CPUs convert instructions (CISC to RISC)


Pipelining

- ARM Cortex-M3 sequence:



Source: ARM

- Idea: Why not fetch the next instruction, while the current one decodes?




Pipelining

- **Timings and definitions**

- Example: Cortex-M3


- **fe: fetch** Read instruction 3 ns¹
 - **de: decode** Decode instruction, read register or memory 4 ns¹
 - **ex: execute** Execute instruction, write back result 5 ns¹




¹timings estimated

Pipelining

– Sequential execution



– Pipelined execution



Pipelining

- **Instructions per second**

- Without pipelining

$$\text{Instructions per second} = \frac{1}{\text{Instruction delay}}$$

- With pipelining

- Pipeline needs to be filled first

- After filling, instructions are executed after every stage

$$\text{Instructions per second} = \frac{1}{\text{Max stage delay}}$$

- **Example Cortex-M3**

- Without pipelining: $1/12\text{ns} = 83,3$ million instructions per second

- With pipelining: $1/5\text{ns} = 200$ million instructions per second

Pipelining

- **Advantages**

- All stages are set to the same execution time
- Massive performance gain
- Simpler hardware at each stage allows for a higher clock rate

- **Disadvantages**

- A blocking stage blocks whole pipeline
- Multiple stages may need to have access to the memory at the same time

- **General**

- Number of execution stages is design decision
- Typically 3 (ARM Cortex M4) – 20 stages (Pentium 4)

Pipelining


– Optimal Pipelining

- All operations here are on registers (single cycle execution)
- In this example it takes 6 clock cycles to execute 6 instructions
- Clock cycles per instruction (CPI) = 1

Cycle		1	2	3	4	5	6	7	8	9
Operation										
<i>ADD</i>	fe	de	ex							
<i>SUB</i>		fe	de	ex						
<i>ORR</i>			fe	de	ex					
<i>AND</i>				fe	de	ex				
<i>ORR</i>					fe	de	ex			
<i>EOR</i>						fe	de	ex		

Pipelining



- **special situation: ldr**
 - In this example it takes 7 clock cycles to execute 6 instructions
 - Read cycle must complete on the bus before ldr instruction can complete
 - Next 2 instructions must wait one pipeline cycle (S = stall)
 - Clock cycles per instruction (CPI) = 1.2



Pipelining

– Control Hazards

- Branch / jump decisions occur in stage 3 (ex)
- Worst case scenario – conditional branch taken:



Pipelining

- **Reduce control hazards**
 - Loop fusion reduces control hazards
 - Simple example for illustration

```
/* Two loops → double number of control hazards*/  
for (i = 0; i < N; i++) {  
    a[i] = 1/b[i] * c[i];  
}  
for (i = 0; i < N; i++) {  
    d[i] = a[i] + c[i];  
}  
  
/* Better performance: Fusion into one loop */  
for (i = 0; i < N; i++) {  
    a[i] = 1/b[i] * c[i];  
    d[i] = a[i] + c[i];  
}
```

2*N Conditional
Branches

N Conditional
Branches

Pipelining

- **Ideas to further improve pipelining:**

- Branch prediction:

- Store last decision made for each conditional branch
 - probability is high that the same decision is taken again

- Instruction prefetch:

- Fetch several instructions in advance
 - better use of the system bus
 - possibility of ‘Out of Order Execution’

- Out of Order Execution:

- If one instruction stalls (e.g. a ldr from slow memory), it might be possible to already execute the next instruction

Pipelining

- **Limits of optimization**

- Complex Optimizations → severe security problems (e.g. Out of Order Execution)
- Instructions (speculatively) executed, that would throw access violations under ‘In Order’ circumstances.
- ‘Meltdown’ and ‘spectre’ attacks: allow a process to access the data of another process.



Parallel Computing

streaming/Vector Processing

- One instruction processes multiple data items simultaneously

Multithreading

- Multiple programs/threads share a single CPU

Multicore Processors

- One processor contains multiple CPU cores

Multiprocessor Systems

- A computer system contains multiple processors

Parallel Computing

- Instructions working on more than one data item

		Data streams	
		Single	Multiple
Instruction streams	Single	SISD : Single processor with one instruction, data stream	SIMD : Data Vectors, all data streams react to one instruction
	Multiple	(no examples)	MIMD : Multiprocessor with SIMD-instructions


Source: David A. Patterson, John L. Hennessy

- **SIMD examples (x86):**
 - MMX (Multimedia Extension), AMD 3D-Now
 - SSE (streaming SIMD Extensions)
 - AVX (Advanced Vector Extensions)

Parallel Computing

- **Multithreading**


- Scheduler: Assigns time slots to programs/threads
- Programs/threads only **seem** to run in parallel (1 CPU)



Remark: More on Multithreading in MC1

Parallel Computing


- **Parallelism on CPU / processor level:**
 - Multicore processor
 - All on one chip
 - Less traffic (cores integrated on one chip)
 - Possibility to share memories on-chip
 - Cheaper
 - Multiprocessor
 - Multiple Chips
 - Longer distances between CPUs
 - More expensive



Source: <https://www.queryhome.com/tech/106320/what-the-difference-between-multicore-and-multiprocessor>

Conclusion

- CPU clock frequency not the only factor
- Different ways to increase processor performance
 - RISC / CISC
 - Pipelining
 - Out-of-order execution
 - Parallel computing
- Today's CPUs combine most of these technologies
- Other components with influence
 - SSD, memory, interfaces, algorithms ...




Source: <https://researchcomputingservices.github.io/parallel-computing/02-speedup-limitations/>

Parallel Computing


- **Multicore processor**

- All on one chip
- Faster running on single program
- Less traffic (cores integrated on one chip)
- No configuration needed
- Cheaper



- **Multiprocessor**

- Multiple chips
- Faster running on multiple programs
- More traffic (distances between chips are higher)
- Configuration needed
- More expensive



Source: <https://www.queryhome.com/tech/106320/what-the-difference-between-multicore-and-multiprocessor>